Claude Team Skills: how it compares

CTS v2.0.0-beta.1 beside the five most-starred agent-workflow frameworks. Competitor columns summarise public documentation and reviews read in November 2026; they describe what each project is known for.

	Claude Team Skills	Superpowers	ECC	BMAD Method	Spec Kit	OpenSpec
GitHub stars	0 (0 forks)	317,000	302,000	55,000	150,000	75,000
Latest release	v2.0.0-beta.1, Nov 2026	v7.0.2, Nov 2026	v2.2.3, Oct 2026	v6.12.4, Oct 2026	v1.4.0, Nov 2026	v2.0.0, Nov 2026
Known for	A 23-role team (engineering, design, security, marketing) run through six gated phases	An engineering methodology: brainstorm, plan, then test-driven work with review	A large ready-made configuration: agents, skills, hooks, security scanning and memory	The full agile lifecycle, from ideation to deployment, with many workflows	Spec-driven development for new work: spec, plan, tasks	Lightweight spec tracking for changes to existing code
How you get it	Claude Code plugin, or copy the markdown; AGENTS.md for other agents	Claude Code plugin; in Anthropic's official marketplace	Plugin for Claude Code, Cursor, Codex and OpenCode	npx installer with selectable modules	Command-line toolkit	Command-line toolkit
Roles beyond engineering	UX, UI, motion, copy, content, marketing, project management	Engineering focused	Engineering focused: planning, review, security, testing agents	Analyst, product and agile roles; creative module	Not role based	Not role based
Hooks and memory	Push gate, session resume and session close hooks; eight project log files	Not covered in the sources reviewed	Hooks across many events; memory persistence	Not covered in the sources reviewed	Not its focus	Not its focus
Fresh-context execution	Implementer subagent per task, then spec review, then quality review	Subagent per task with spec-compliance and code-quality review	Specialised agents per activity	Party mode: several personas discuss in one session	Not its focus	Not its focus
Specs for existing code	Change-spec workflow: baseline, proposal, delta, spec check, merge	Plans per task	Not its focus	Planning documents per track	Aimed at new work	Its core purpose
Scaling the ceremony	Three tracks: Quick, Standard, Full	One workflow	Not covered in the sources reviewed	Tracks from quick flow to enterprise	One workflow	Deliberately lightweight
Common criticism in reviews	No users yet; small eval suite (see page 2)	Overhead and latency from mandatory steps	No independent critical review found	Setup effort; heavy for teams under five	Heavy for small changes	Less guidance for building new systems

What sets CTS apart

What CTS does	How it is backed	Where to look
Blocks a push until verification is logged	A hook stops **git push** unless the latest VERIFY entry is a pass, is newer than the last SHIP, and logged coverage is at least 80%. The same script works as a plain git pre-push hook under any agent.	hooks/push-gate.sh, 13 behaviour tests
Covers the whole product, not only the code	23 specialists, including UX, UI, motion, security, DBA, copy, content and marketing, loaded one at a time. This is the clearest gap in the engineering-focused alternatives.	skills/
Remembers the project between sessions	Eight append-only log files. A hook loads the last session summary at start and closes sessions that ended without one.	hooks/session-start.sh, session-end.sh
Scales effort to the request	Quick answers come straight from one specialist; Standard and Full add gates, documents and subagents. All six phases are logged on every track.	skills/orchestrator/SKILL.md
Reviews with fresh eyes	Code tasks go to an implementer, then a spec reviewer, then a quality reviewer, none of which saw the conversation. Visual work goes through three independent design critics.	agents/, skills/design-loop/
Publishes its own test results	Behavioural evals run with Claude Code's built-in eval command; results and known gaps are in the repo.	evals/README.md

Known limits

• **No adoption yet.** On 13 November 2026 the repository had 0 stars, 0 forks and 5 unique cloners in 14 days. Every project in the table has between 55,000 and 317,000 stars.

• **Hooks, memory and subagents are not unique.** ECC ships hooks and memory, and Superpowers ships per-task subagents with two-stage review. What CTS combines them with is the non-engineering team and a gate on the push itself.

• **Evals are small.** Four core cases pass, one run each. A fifth check, that the orchestrator loads when a user demands to ship untested code, failed in 6 of 6 runs and was removed from scoring: Claude holds the gate, but on its own.

• **The review loop has not been run end to end.** Its eval case is written but needs a sandbox (WSL2 or Linux) that was not available when this was published.

• **The push gate checks the log, not the tests.** It catches a skipped VERIFY phase. It cannot tell whether a logged pass was earned.

• **Hooks and subagents need the plugin install,** and AGENTS.md is untested outside Claude Code.

Sources for the competitor columns

GitHub API, 13 November 2026 (stars, forks, releases) · rywalker.com/research/agentic-skills-frameworks · dailyaiworld.com/blogs/ai-agent-skills-frameworks-comparison-2026 · mintlify.com/bmad-code-org/BMAD-METHOD · mintlify.com/obra/superpowers/concepts/workflow · kunalganglani.com/blog/superpowers-agentic-skills-framework · augmentcode.com/learn/everything-claude-code-github · mintlify.com/affaan-m/everything-claude-code · dabase.com/blog/2026/sdd-framework-comparison · codemyspec.com/blog/openspec-vs-spec-kit

Star counts are rounded and come from the GitHub API on the date shown. Two other large projects are not in the table: anthropics/skills (194,000 stars, a skill catalogue with no workflow) and get-shit-done (64,000 stars, no commits since May 2026).